

Markov Chain Monte Carlo Methods

Nikolas Tsigkas

May 2021

1 Gibbs Sampler for a Normal Model

The given dataset contains daily precipitation data, measured in $\frac{1}{100}$ of an inch, from Snoqualmie Falls, Washington from 1948 to 1983. The proposed model for the data is: $\ln(Y_1), \dots, \ln(Y_n) | \mu, \sigma \stackrel{iid}{\sim} N(\mu, \sigma^2)$, where both μ and σ are unknown in the Bayesian sense. They are in turn modelled as: $\mu \sim N(\mu_0, \tau_0^2)$ and $\sigma^2 \sim \text{Scale-inv-}\chi^2(\nu_0, \sigma_0^2)$. The joint distribution of (μ, σ^2) can be sampled by the Gibbs Sampling Monte Carlo algorithm. The algorithm works as follows:

Algorithm 1: Gibbs Sampling

```
initialise the parameters  $\theta_1^{(0)}, \dots, \theta_p^{(0)}$ ;  
for  $i = 1, \dots, N$  do  
  for  $j = 1, \dots, p$  do  
    | Draw  $\theta_j^{(i)}$  from the conditional posterior  $p(\theta_j | \theta_1^i, \dots, \theta_{j-1}^i, \theta_{j+1}^{(i-1)}, \dots, \theta_p^{(i-1)})$   
  end  
end  
Result:  $(\theta_1^{(1)}, \dots, \theta_p^{(1)}), \dots, (\theta_1^{(N)}, \dots, \theta_p^{(N)})$ 
```

The Gibbs Sampling method is a convenient method for when the full conditional distributions are known, which is the case here, as it can be shown that:

$$\begin{aligned} w &= \frac{n/\sigma^2}{n/\sigma^2 + 1/\tau_0^2} \\ \mu_n &= w\bar{y} + (1-w)\mu_0 \\ \frac{1}{\tau_n^2} &= \frac{1}{n/\sigma^2 + 1/\tau_0^2} \\ \mu | \sigma, \mathbf{y} &\sim N(\mu_n, \tau_n) \\ \sigma^2 | \mu, \mathbf{y} &\sim \text{Scale-Inv-}\chi^2(\nu_0 + n, \frac{\nu_0 \sigma_0^2 + \sum_{i=1}^n (y_i - \mu)^2}{\nu_0 + n}) \end{aligned}$$

The algorithm was performed using the prior parameters: $\mu_0 = 3.7$, $\tau_0 = 100$, $\nu_0 = 1000$ and $\sigma_0 = 1$, as well as using the initial values $\mu^{(0)} = 1$, $\sigma^{(0)} = 1$. The resulting scatter plot of 5000 random draws from the Gibbs-sampler are shown in figure 1.

One possible problem of the Gibbs sampler, as with all Markov Chain Monte Carlo methods, is that the samples can be heavily autocorrelated, resulting in a slower convergence towards the mean. This is due to the fact that the variance of the sample mean of a parameter θ becomes: $\mathbf{V}[\frac{1}{N} \sum_{i=1}^N \theta^{(i)}] = \frac{\sigma^2}{N} [1 + 2 \sum_{i=1}^N \text{acf}(i)]$, where $\text{acf}(i)$ denotes the autocorrelation at lag i . The term inside the parenthesis is often referred to as the inefficiency factor (IF), and can give a good description of how fast the simulations converges. From the simulations, this quantity was estimated to be $\text{IF} = 5.168$.

Another way to examine the convergence is visually, by plotting the Markov Chain trajectories. Figure 2 shows that the algorithm draws values from what appears to be a stationary distribution.

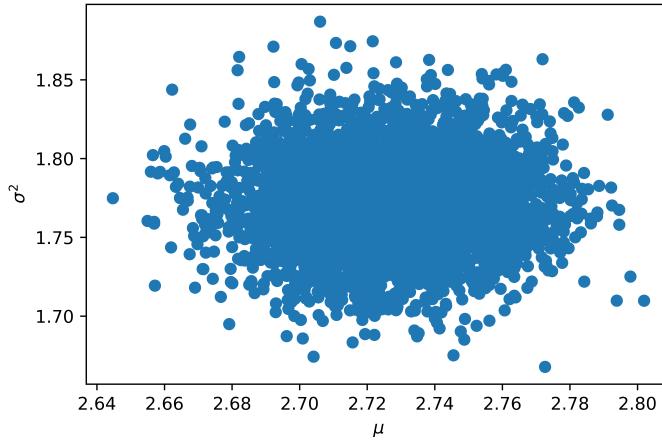


Figure 1: Scatter plot over the simulated μ and σ^2 values

Finally, the posterior predictive density can be evaluated using the Gibbs draws as:

$$p(y_{n+1}|\mathbf{y}) = \iint_{\Theta} p(y_{n+1}|\mathbf{y}, \theta)p(\theta|\mathbf{y}) = \mathbf{E}_{\theta|\mathbf{y}}[p(y_{n+1}|\theta)] \approx \frac{1}{N} \sum_1^N p(y_{n+1}|\theta^{(i)})$$

Where Θ is the domain space of the parameters, N is the number of Gibbs draws and $\theta^{(i)}$ is the vector of Gibbs draws, i.e. $(\mu^{(i)}, (\sigma^2)^{(i)})$, at iteration i . By using the Gibbs draws, one can then draw from the posterior predictive density by drawing random samples from a $N(\mu^{(i)}, \sigma^{(i)})$ distribution. This resulting distribution can then be compared to the observed data. Figure 3 shows the log of the observed data in a histogram together with the posterior predictive draws. The results are not too promising, as there is quite some deviation between the data and the pdf. For better results, one could examine for instance normal mixture distributions, with which could potentially capture the skewness which is apparent in the plot. Drawing from mixture densities is also fairly convenient with the Gibbs sampling algorithm.

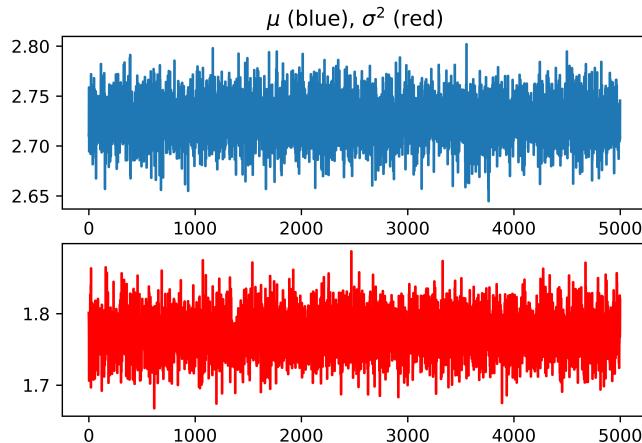


Figure 2: The Markov Chain trajectories for the sampled parameters

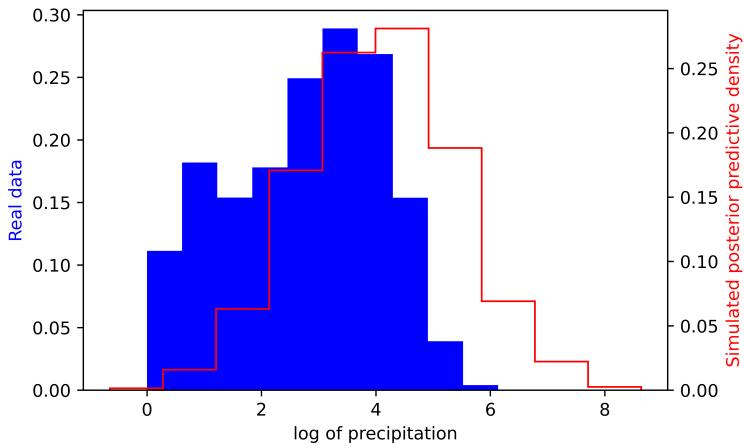


Figure 3: The real data compared to the estimated posterior predictive density

```

1 """ Some useful packages"""
2 import pandas as pd
3 import numpy as np
4 from scipy import stats as st
5 import matplotlib.pyplot as plt
6
7 np.random.seed(1111) # set seed for reproducability
8
9 """ Import and transform the data """
10 data = pd.read_csv('rainfall.dat', header = None)
11 log_y = np.log(data.values)
12
13 """ Helper functions"""
14 # Draws 1 sample from the Scale-Inverse-chi2 distribution
15 def draw_chi2inv(df, scale):
16     x = st.chi2.rvs(df-1, scale, size=1)
17     sigma_sq = (df-1)*scale/x
18     return sigma_sq
19
20 # Returns the parameters for the posterior of mu
21 def mu_post_params(pri_params, sigma, data):
22     [mu_0, tau_0, nu_0, sigma_0] = pri_params
23     n = len(data)
24     w = (n/sigma**2)/(n/sigma**2 + tau_0**(-2))
25     mu_n = w*np.mean(data)+(1-w)*mu_0
26     tau_n = (1/(n/sigma**2 + 1/tau_0**2))**0.5
27     return [mu_n, tau_n] # Mean and sd for the normal distribution
28
29 # Returns the parameters for the posterior of sigma^2
30 def sigma_post_params(pri_params, mu, data):
31     [mu_0, tau_0, nu_0, sigma_0] = pri_params
32     n = len(data)
33     nu_n = nu_0 + n
34     scale = (nu_0*sigma_0**2 + np.sum((data-mu)**2) )/(nu_n)

```

```

35     return [nu_n, scale] #DoF and scale parameter for the Scale-Inv-chi2
36
37 # Draws nDraws Gibbs-samples from the posterior
38 def Gibbs(pri_params, post_params1, post_params2, dist1, dist2, iVals, nDraws,
39            data):
40     var1 = np.zeros(nDraws+1)
41     var2 = np.zeros_like(var1)
42     # Initialize the Gibbs algorithm with the initial values
43     [var1[0], var2[0]] = iVals
44     # Sample iteratively from the conditional marginal distributions
45     for i in range(1,nDraws+1):
46         [mu_n, tau_n] = post_params1(pri_params, var2[i-1], data)
47         var1[i] = dist1(mu_n, tau_n)
48         [nu_n, scale] = post_params2(pri_params, var1[i], data)
49         var2[i] = dist2(nu_n, scale)
50     return [var1[1:], var2[1:]]
51
52 # Sample autocorrelation
53 def acf(x):
54     n = len(x)
55     autocorr = []
56     for i in range(1, len(x)-1):
57         autocorr.append(np.corrcoef(x[:n-i],x[i:])[0,1])
58     return np.array(autocorr)
59
60 """ Prior parameters """
61 # pri_params = mu_0, tau_0, nu_0, sigma_0
62 pri_params = np.array([3.7, 100.0, 1000.0, 1.0])
63 iVals = np.ones(2)
64
65 [sample_mu, sample_sigma2] = Gibbs(pri_params,
66                                     mu_post_params, sigma_post_params,
67                                     st.norm.rvs, draw_chi2inv,
68                                     iVals,
69                                     5000,
70                                     log_y
71                                     )
72
73 # Calculate the inefficiency factor
74 autocorr = acf(sample_sigma2)
75 IF = 1 + 2*np.sum(autocorr)
76
77 # Scatterplot of the simulated draws
78 plt.scatter(sample_mu, sample_sigma2)
79 plt.xlabel(r"\mu")
80 plt.ylabel(r"\sigma^2")
81 plt.show()
82
83 # Trajectories of the two Markov chains
84 fig, axs = plt.subplots(2,1)
85 axs[0].plot(sample_mu)
86 axs[1].plot(sample_sigma2, color = "r")
87
88 axs[0].set_title(r"\mu (blue), \sigma^2 (red)")

```

```

88 plt.savefig("Gibbs_Convergence.png", dpi = 800)
89 plt.show
90
91 # Compare the empirical data to the posterior predictive distribution
92 # from the MCMC simulation
93 y_pred = np.zeros((5000,1))
94
95 for i in range(len(y_pred)):
96     y_pred[i] = sample_mu[i] + np.sqrt(sample_sigma2[i])*st.norm.rvs(1)
97
98 fig, ax1 = plt.subplots()
99 ax2 = ax1.twinx()
100 ax1.hist(log_y, density = True, color = "b")
101 ax2.hist(y_pred, color="r", histtype = "step", density = True)
102
103 ax1.set_xlabel('log of precipitation')
104 ax1.set_ylabel('Real data', color='b')
105 ax2.set_ylabel('Simulated posterior predictive density', color='r')
106 plt.show()

```

2 Random Walk Metropolis for Poisson Regression

The given dataset contains data on from coin auctions on eBay. The goal is to fit a model that explains the number of bids that an auction receives, given the covariates:

- If the seller is a "Power Seller", meaning they put up a high volume of items for auction (0,1)
- If the seller is verified by eBay (0,1)
- If the sold coin was packaged in an unopened envelope (0,1)
- If the item has a minor or major defect (0,1 for each type of defect)
- If the seller has recorded a lot of negative feedback (0,1)
- The log of the book value of the coin, as deemed by experts
- The ratio between the minimal bidding price and the book value

Formally, the model is:

$$Y_i | \mathbf{x}_i, \beta \sim \text{Po}[\exp(\mathbf{x}_i^T \beta)]$$

Where y_i is the recorded number of bids, and \mathbf{x}_i is a vector containing the covariates of the auction, with an intercept term added. Initially, the maximum likelihood estimate of the regression coefficients β can be explored, for instance using the STATSMODEL package in Python. The results from this are shown in the code snippet below.

1	print(poisson_fit.summary())					
Generalized Linear Model Regression Results						
<hr/>						
4	Dep. Variable:	y	No. Observations:	1000		
5	Model:	GLM	Df Residuals:	991		
6	Model Family:	Poisson	Df Model:	8		
7	Link Function:	log	Scale:	1.0000		
8	Method:	IRLS	Log-Likelihood:	-1796.1		
9	Date:	Thu, 13 May 2021	Deviance:	867.47		
10	Time:	19:00:29	Pearson chi2:	752.		
11	No. Iterations:	5				
12	Covariance Type:	nonrobust				
<hr/>						
14	coef	std err	z	P> z	[0.025	0.975]
15						
16	const	1.0724	0.031	34.848	0.000	1.012
17	x1	-0.0205	0.037	-0.558	0.577	-0.093
18	x2	-0.3945	0.092	-4.268	0.000	-0.576
19	x3	0.4438	0.051	8.778	0.000	0.345
20	x4	-0.0522	0.060	-0.867	0.386	-0.170
21	x5	-0.2209	0.091	-2.416	0.016	-0.400
22	x6	0.0707	0.056	1.255	0.210	-0.040
23	x7	-0.1207	0.029	-4.166	0.000	-0.177
24	x8	-1.8941	0.071	-26.588	0.000	-2.034
25						

One can see that the variables x_1 , x_4 and x_6 , corresponding to the "Power seller", minor damage and negative feedback parameters, are not statistically significant on a 5% level from the resulting p-values.

Furthermore, one can also draw simulated values for these parameters using for instance the Metropolis Random Walk algorithm. The algorithm works by performing the following steps:

Algorithm 2: Random Walk Metropolis

```

for  $i = 1, \dots, N$  do
    Draw a sample proposal  $\theta^{(p)}$  from the conditional probability  $\Theta^{(p)} | \theta^{(i-1)} \sim N(\theta^{(i-1)}, c\Sigma)$ ;
    Compute the acceptance probability  $\alpha_i = \min(1, \frac{p[\theta^{(p)} | \mathbf{y}]}{p[\theta^{(i-1)} | \mathbf{y}]})$ ;
    Draw a sample  $x_i$  from  $X_i \sim \text{Bern}(\alpha_i)$ ;
    if  $x_i = 1$  then
        |  $\theta^{(i)} \leftarrow \theta^{(p)}$ 
    else
        |  $\theta^{(i)} \leftarrow \theta^{(i-1)}$ 
    end
end

```

Two possible choices for the Σ parameter is the identity matrix, or the negative inverse hessian of the log-posterior, evaluated at the posterior mode. The latter was found through numerical optimisation of the log-posterior pdf, using the Broyden-Fletcher-Goldfarb-Shanno algorithm. The prior used for the β parameters was the so called Zellner g-prior:

$$\beta \sim N(0, 100[\mathbf{X}^T \mathbf{X}]^{-1})$$

This gives a log posterior proportional to:

$$\ln(p(\beta | \mathbf{y})) \propto \sum_{i=1}^n \ln[p(y_i | \beta)] + \ln[p(\beta)] = \sum_{i=1}^n [-\lambda_i + y_i \ln(\lambda_i) - \sum_{j=1}^{y_i} \ln(j)] + \phi_{\mathbf{0}, 100\mathbf{X}^T \mathbf{X}^{-1}}(\beta)$$

Where $\lambda_i = \exp(\mathbf{x}_i^T \beta)$ and $\phi_{\mu, \Sigma}$ is the pdf of the normal distribution with mean vector and covariance matrix μ and Σ respectively. It can be noted that, for the optimisation, the inner sum in the first term is redundant and can thus be omitted, as it is independent of β . The resulting posterior mode and negative inverse hessian of the log posterior was:

$$\tilde{\beta} = [1.0698 \quad -0.0205 \quad -0.3930 \quad 0.4436 \quad -0.0525 \quad -0.2212 \quad 0.0707 \quad -0.1202 \quad -1.8920]^T$$

$$-H_{\ln p}^{-1}(\tilde{\beta}) = \begin{bmatrix} 0.0011 & -0.0007 & 0 & -0.0004 & -0.0003 & 0.0002 & -0.0004 & 0 & 0.0012 \\ -0.0007 & 0.0013 & -0.0002 & -0.0002 & 0.0006 & 0.0008 & 0.0004 & 0 & -0.0002 \\ 0 & -0.0002 & 0.0027 & 0 & 0.0007 & 0.0024 & 0.0007 & -0.0014 & -0.0009 \\ -0.0004 & -0.0002 & 0 & 0.0021 & -0.0005 & -0.0008 & -0.0002 & 0 & -0.0008 \\ -0.0003 & 0.0006 & 0.0008 & -0.0005 & 0.0035 & 0.0038 & 0.0002 & -0.0004 & -0.0007 \\ 0.0002 & 0.0008 & 0.0024 & -0.0008 & 0.0038 & 0.0223 & 0.0027 & -0.0012 & 0.0010 \\ -0.0004 & 0.0004 & 0.0007 & -0.0002 & 0.0002 & 0.0027 & 0.0029 & -0.0004 & -0.0003 \\ 0 & 0 & -0.0014 & 0 & -0.0004 & -0.0012 & -0.0004 & 0.0007 & 0.0005 \\ 0.0012 & -0.0002 & -0.0009 & -0.0008 & -0.0007 & 0.0010 & -0.0003 & 0.0005 & 0.0034 \end{bmatrix}$$

Furthermore, the ratio of posterior probabilities, needed to determine α in the Metropolis algorithm can be calculated using the log-posterior, as: $\exp[\ln(p(\beta^i | \mathbf{y})) - \ln(p(\beta^{i-1} | \mathbf{y}))]$, where once again the terms independent with regards to β can be neglected. This approach is also more numerically sound, as the originally defined ratio can lead to division by very small numbers.

The final terms needed to perform the Monte Carlo algorithm are the initial value of β , which was set to the zero-vector, and the constant c . This parameter can be used to tune the sampling to get satisfactory results. One rule of thumb is to have an average acceptance probability within the region of 25 – 30%. With $c = 0.6$, the average acceptance was 27.1%. The resulting marginal trajectories for each element in the β -vector can also be examined to see whether the algorithm has converged. Figure 4 shows that the samples of all the elements follow a stationary process after around 500 iterations.

From these random draws, the sample mean can be evaluated. By eliminating the first 500 random draws, the so called "burn-in" period, the sample mean was estimated to:

$$\hat{\beta} = [1.0612 \quad -0.0097 \quad 0.0512 \quad 0.3736 \quad -0.0517 \quad -0.2173 \quad 0.057 \quad -0.0256 \quad -1.770]^T$$

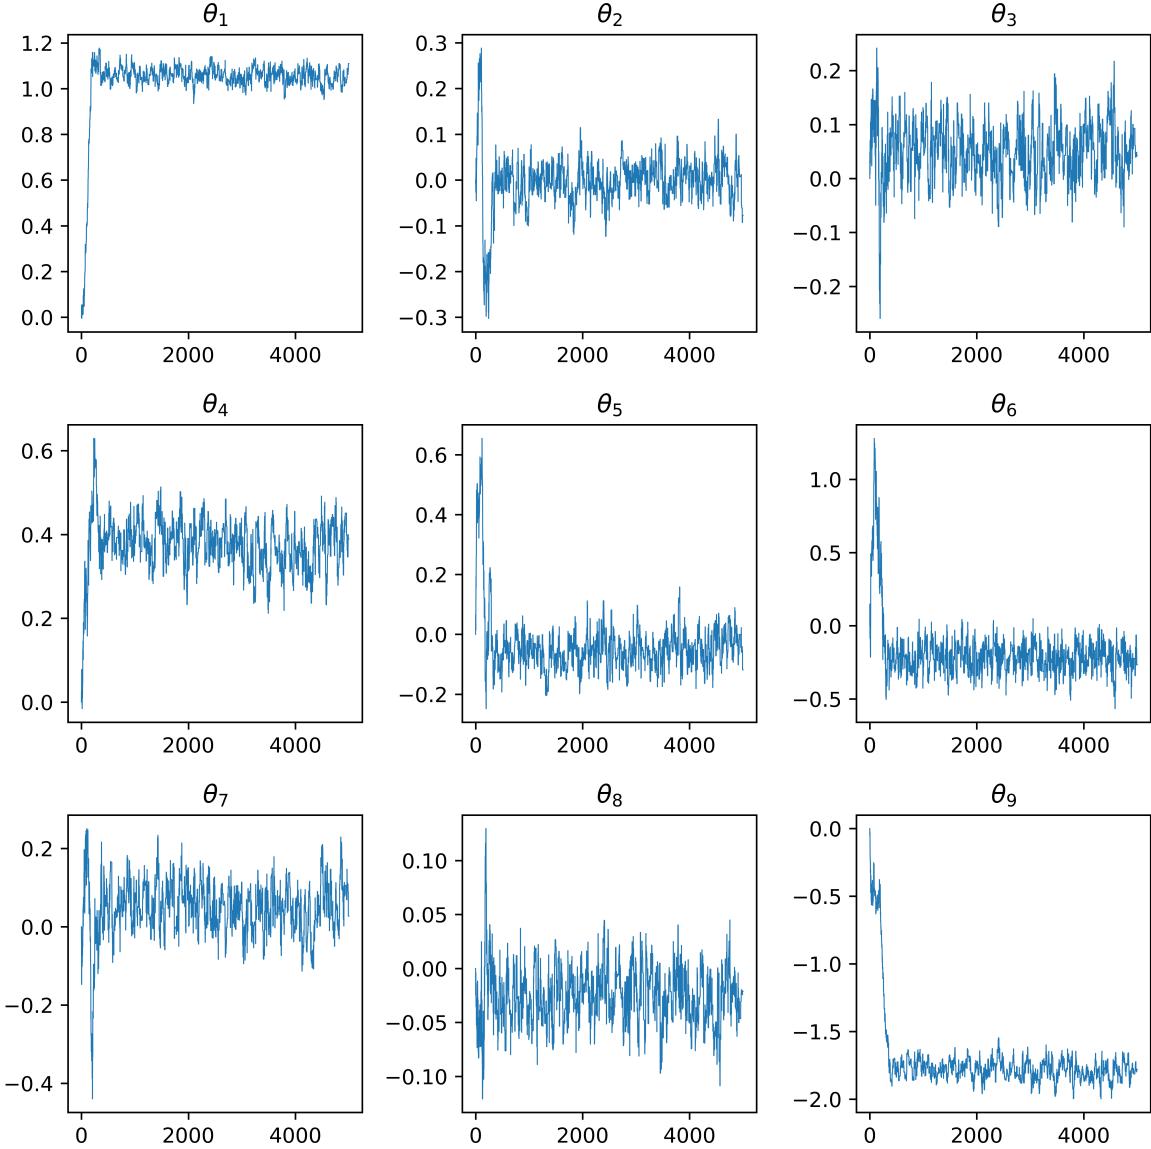


Figure 4: Trajectories of the regression parameters for 5000 simulations

Using the draws from the stationary distribution, one can examine the distribution of a new, unobserved sample. For instance the posterior predictive distribution of an auction with:

- A power seller with a verified ID
- A sealed product with a major defect, but no minor one
- No major negative feedback of the bidder
- A log book value of 1
- A minimum bid to book value of 0.7

Which corresponds to the data vector:

$$\mathbf{x}_s = [1 \ 1 \ 1 \ 1 \ 0 \ 1 \ 0 \ 1 \ 0.7]^T$$

Can be evaluated using the same procedure as with the Gibbs sampling algorithm, i.e. drawing random values from a $Po[\exp(\mathbf{x}_s^T \beta^{(i)})]$ distribution for the different $\beta^{(i)}$ draws produced from the Metropolis Random Walk. The histogram

from these draws, together with the pmf corresponding to the maximum likelihood estimate of the statsmodel package is shown in figure 5. From these draws, the probability that this auction will have no bids, $P[Y_s = 0|\beta]$ was also evaluated to be approximately 18.69%

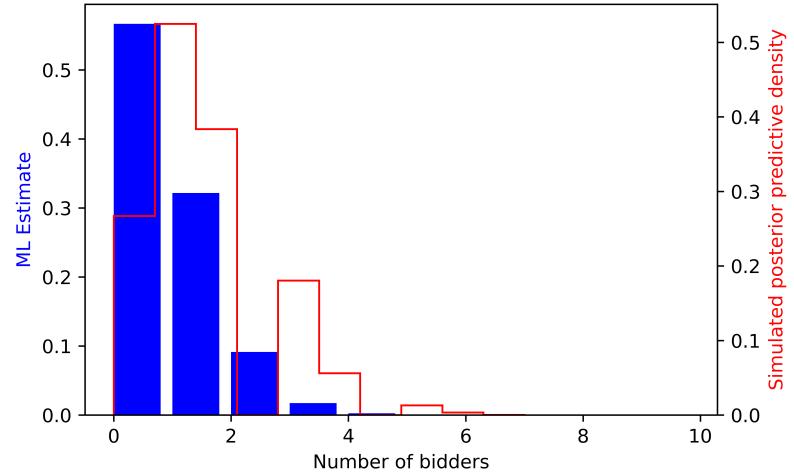


Figure 5: The Maximum Likelihood estimate compared to the posterior predictive density

```

1  """ Some useful packages"""
2 import pandas as pd
3 import numpy as np
4 from scipy import stats as st
5 import matplotlib.pyplot as plt
6 import statsmodels.api as sm
7 from scipy.optimize import minimize as fmin
8
9 np.random.seed(1111) # set seed for reproducability
10
11 """ Helper functions """
12 # An expression proportional to the log of the posterior, minus the terms
13 # that are not dependent on beta,
14 # to be used for the numerical optimization and the MRW
15 def log_posterior(beta, X, y):
16     regr = X@beta
17     logL = np.sum(-np.exp(regr) + y*regr)
18     prior = np.log(st.multivariate_normal.pdf(beta, mean=np.zeros_like(beta), cov =
19         100*np.linalg.inv(X.T@X)))
20     return (logL + prior)
21
22 class RWM_res:
23     def __init__(self, init, samples):
24         self.initial = init
25         self.iterations = samples
26         self.theta = np.zeros((len(init), samples+1))
27         self.theta[:,0] = self.initial # Initialize
28         self.acc_p = 0.0 # Keep track of the acceptance probability
29     def accept(self, theta_p, it):
30         self.acc_p += 1.0/(self.theta.shape[1])
31         self.theta[:,it] = theta_p
32     def reject(self, it):
33         self.theta[:,it] = self.theta[:,it-1]
34     def draw(self, it, Cov):
35         return st.multivariate_normal.rvs(mean=self.theta[:,it-1], cov=Cov)
36     def hat(self, burn_in):
37         return np.mean(self.theta[:,burn_in:], axis=1)
38     def plot(self, x, y):
39         order = np.reshape(np.arange(0,x*y),(x,y))
40         fig, axs = plt.subplots(x,y, figsize = (8,8))
41         for i in range(x):
42             for j in range(y):
43                 axs[i,j].plot(self.theta[order[i,j],:], linewidth=0.5)
44                 axs[i,j].set_title(r"$\theta_{%i}$" %(order[i,j]+1))
45         fig.tight_layout()
46         plt.show()
47     def iterate(self, log_post_fun, X, y, Cov, c):
48         for i in range(1,self.iterations+1):
49             theta_p = self.draw(i, c*Cov)
50             # Calculate the probability of acceptance
51             post_ratio = np.exp(log_post_fun(theta_p, X, y) - log_post_fun(self.theta
52                 [:,i-1], X, y))
53             # Draw a random sample from a Bern(alfa) distribution
54             # To determine if the iteration should be accepted or not
55             alfa = np.nanmin([1, post_ratio])
56             if st.bernoulli.rvs(alfa):
57                 self.accept(theta_p, i)
58             else:

```

```

57         self.reject(i)
58
59 """ Import and transform the data """
60 data = pd.read_csv('eBayNumberOfBidderData.csv', sep = ";")
61
62 y = data.nBids.values # Response variable
63 X = data.iloc[:, 1:].values # Covariates
64
65 """Fit a generalized linear model using the statsmodel package"""
66 exog, endog = X, y
67 model = sm.GLM(endog, exog,
68                 family = sm.families.Poisson(link=sm.families.links.log))
69
70 poisson_fit = model.fit()
71 print(poisson_fit.summary())
72
73 """Find the normal approximation of the posterior, to get the inverse hessian"""
74 # Define the log posterior as a function of beta only, for the optimization
75 # Also, multiply with -1, as we seek the maximum
76 # and the algorithm performs minimization
77 def obj_fun(beta):
78     return -1*log_posterior(beta, X, y)
79
80 # Find the max of the log posterior using a quasi-Newton method
81 # Which also returns a numerical approximation of H^-1
82 x0 = np.zeros(9) # Initial guess
83 opt_res = fmin(obj_fun, x0, method="BFGS", options={"disp" : True})
84
85 # Numerical approximation of -H^-1
86 H_inv = opt_res["hess_inv"]
87 # The mode, i.e. the beta vector at the optimum
88 post_mode = opt_res['x']
89
90 """Draw from the actual posterior,
91 using the Random Walk Metropolis algorithm"""
92 RWM_results = RWM_res(np.zeros(9), 5000) # Initialize
93 RWM_results.iterate(log_posterior, X, y, H_inv, 0.6)
94
95 RWM_results.acc_p # Check the average acceptance probability
96
97 # Plot the marginal trajectories to check convergence visually
98 RWM_results.plot(3,3)
99 plt.savefig("MCMC_convergence.png", dpi=800)
100
101 # Compute the sample mean, eliminating the first 500 iterations
102 beta_hat = RWM_results.hat(500)
103
104 # Using the iterations to sample from the posterior predictive density
105 beta_draws = RWM_results.theta[:, 501:]
106
107 # For the following auction:
108 """
109 The x-vector for an auction with:
110     A power seller,
111     with verified ID.
112     A sealed product with
113     a major defect
114     No major negative feedback
115     A log book-value of 1

```

```

116     A minimum bid of 0.7
117 """
118 x_sample = np.array([1,1,1,1,0,1,0,1,0.7])
119
120 y_pred = np.zeros(4500)
121 for i in range(len(y_pred)):
122     lamda = np.exp(x_sample@beta_draws[:, i])
123     y_pred[i] = st.poisson.rvs(1, lamda)
124
125 # The posterior distribution of this sample
126 lamda_MLE = np.exp(x_sample@poisson_fit.params)
127 y_MLE = st.poisson.pmf(np.arange(0,10), lamda_MLE)
128
129 fig, ax1 = plt.subplots()
130 ax2 = ax1.twinx()
131 ax1.bar(np.arange(0,10), y_MLE, color = "b", align = "edge")
132 ax2.hist(y_pred, color="r", histtype = "step", density = True, align = "mid")
133
134 ax1.set_xlabel('Number of bidders')
135 ax1.set_ylabel('ML Estimate', color='b')
136 ax2.set_ylabel('Simulated posterior predictive density', color='r')
137 plt.show()
138
139 # The probability that the sample gets no bids
140 p0 = np.sum(y_pred == 0)/len(y_pred)

```

3 Hamiltonian Monte Carlo for Time Series Analysis

The aim is to use Hamiltonian Monte Carlo, through the PyStan package, to evaluate parameters of an $AR(1)$ process $\{X_t\}_{t=0}^T$, defined in the following way:

$$X_t = \mu + \phi(X_{t-1} - \mu) + \epsilon_t \quad \epsilon_t \stackrel{iid}{\sim} N(0, \sigma) \quad X_0 = \mu$$

First, two synthetic datasets from this process are generated; $\{X_t\}_{t=0}^T$, which has parameters $\mu = 20$, $\phi = 0.3$, $\sigma = 2$ and $T = 200$, and $\{Y_t\}_{t=0}^T$, which has parameters $\phi = 0.9$ and μ and σ are the same as for X_t . The trajectories for the two processes are shown in figure 6.

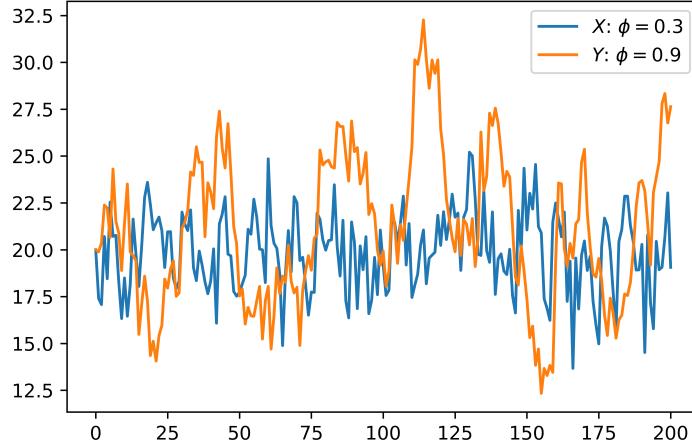


Figure 6: The trajectories of the two $AR(1)$ processes

Next, the parameters μ , σ and ϕ are treated as unknown parameters in the Stan model, using the priors:

$$p(\phi) \propto \begin{cases} 1, & -1 \leq \phi \leq 1 \\ 0 & \text{otherwise} \end{cases} \quad p(\sigma) \propto \begin{cases} c, & 0 < \sigma \\ 0, & \text{otherwise} \end{cases} \quad p(\mu) \propto c$$

As $|\phi| > 1$ will lead to an unstable process (in general, the impact of ϕ is the autocorrelation between two subsequent samples), and σ must be positive. The way this is defined in Stan is showed in the code snippet below.

```

1 data {
2     int<lower=0> T;
3     vector [T] x;
4 }
5 parameters {
6     real<lower=-1, upper=1> phi;
7     real mu;
8     real<lower=0> sigma;
9 }
10 model {
11     for (t in 2:T)
12         x[t] ~ normal(mu + phi * (x[t-1] - mu), sigma);
13 }
```

Using this model, 5000 iterations per chain and four chains where performed for each process, using the "No U-Turn Sampling" algorithm in Stan. For the first, a warm-up period of 1000 was used while, for the second, the warm-up period was 2500, as a way of increasing the robustness of the model. The results from the Monte Carlo simulations are shown below, in terms of the model output shown below, and the pairwise plots of the posterior parameters shown in figure 7.

```

1 x_fit
2
3 Inference for Stan model: anon_model_bed26d3d7e672f0db82b8a542ddb1573 .
4 chains, each with iter=5000; warmup=1000; thin=1;
5 post-warmup draws per chain=4000, total post-warmup draws=16000.
6
7      mean se_mean     sd 2.5%   25%   50%   75% 97.5% n_eff Rhat
8 phi    0.36  5.7e-4  0.07  0.22  0.31  0.36  0.4   0.49 14043  1.0
9 mu     19.88 2.1e-3  0.23 19.43 19.73 19.88 20.03 20.33 11258  1.0
10 sigma  2.01  8.4e-4  0.1   1.83  1.94  2.01  2.08  2.23 14647  1.0
11 lp__ -239.4   0.01  1.27 -242.6 -239.9 -239.0 -238.4 -237.9  7864  1.0
12
13 Samples were drawn using NUTS at Sat May 15 15:55:03 2021.
14 For each parameter, n_eff is a crude measure of effective sample size,
15 and Rhat is the potential scale reduction factor on split chains (at
16 convergence, Rhat=1).

```

```

1 y_fit
2
3 Inference for Stan model: anon_model_06adc118d3f2b3d73f4283204fec27dc .
4 chains, each with iter=5000; warmup=2500; thin=1;
5 post-warmup draws per chain=2500, total post-warmup draws=10000.
6
7      mean se_mean     sd 2.5%   25%   50%   75% 97.5% n_eff Rhat
8 phi    0.91  1.8e-3  0.04  0.84  0.89  0.91  0.94  1.0   475  1.01
9 mu     29.78  5.55  74.35 16.6  20.54 21.57 22.73 46.14  179  1.02
10 sigma  1.87  1.7e-3  0.1   1.69  1.81  1.87  1.93  2.07 3250  1.0
11 lp__ -226.6   0.17  2.35 -233.7 -227.3 -225.9 -225.1 -224.4  187  1.03
12
13 Samples were drawn using NUTS at Sat May 15 16:27:01 2021.
14 For each parameter, n_eff is a crude measure of effective sample size,
15 and Rhat is the potential scale reduction factor on split chains (at
16 convergence, Rhat=1).

```

From this one can extract the following characteristics of the sample posterior:

	X			Y		
	$\hat{\mu}$	$CI_{0.95}$	n_{eff}	$\hat{\mu}$	$CI_{0.95}$	n_{eff}
ϕ	0.36	0.22	0.49	14043	0.91	0.84
μ	19.88	19.43	20.33	11258	29.78	16.6
σ	2.01	1.83	2.23	14647	1.87	1.69
					2.07	3250

Table 1: Parameter data estimated from the HMC simulations

From the results summary, as well as the plots, it can clearly be seen that the samples from X can be used to accurately describe the parameters, while they fail to do so for Y . One possible reason for this is that the subsequent samples of Y have much heavier autocorrelation, which reduces the efficiency of the sampling. This is noted when looking at the number of efficient draws from Y , which is extremely low. Another possible issue was the selected prior for ϕ between -1 and 1, as this prohibits sampling much higher than the true value of ϕ .

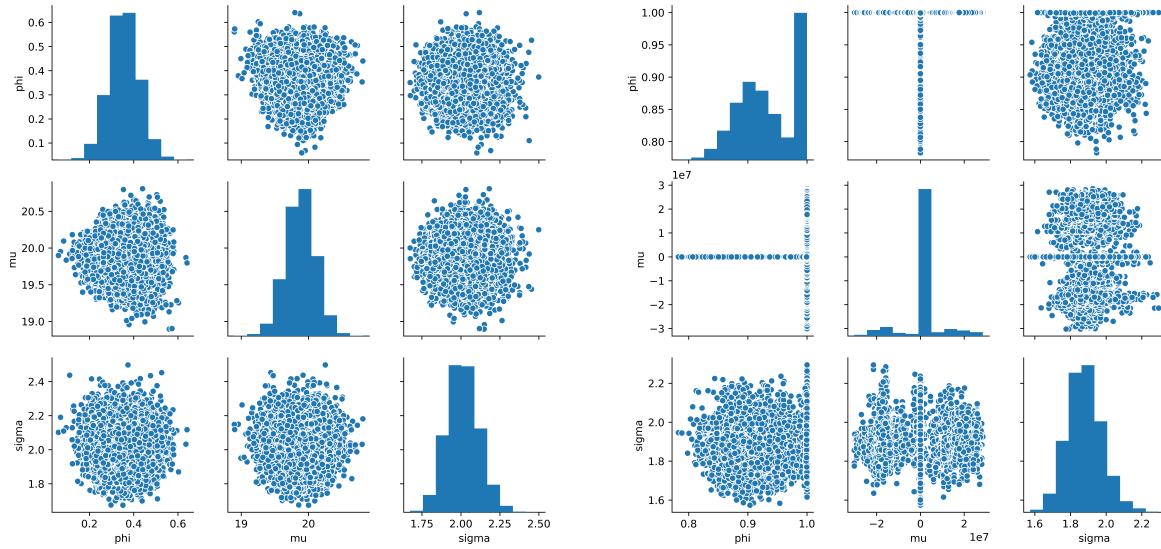


Figure 7: Pairs plots for the parameters of the two processes (left: X , right: Y)

```

1 """ Some useful packages """
2 import numpy as np
3 from scipy import stats as st
4 import matplotlib.pyplot as plt
5 import pystan as stan
6 import seaborn as sea
7
8 np.random.seed(1111) # Set seed for reproducibility
9
10 """ Functions that draws samples from an AR(1)-process """
11 def AR1(mu, phi, sigma, T, x0):
12     x = np.zeros(T+1)
13     x[0] = x0
14     for t in range(1, T+1):
15         x[t] = st.norm.rvs(loc=mu+phi*(x[t-1]-mu), scale=sigma)
16     return x
17
18 """ Explore the impact of phi on x """
19 x = AR1(20, 0.3, 2, 200, 20)
20 y = AR1(20, 0.9, 2, 200, 20)
21
22 plt.plot(x)
23 plt.plot(y)
24 plt.legend([r"$X$: $\phi=0.3$","$Y$: $\phi=0.9$"])
25
26 """ Define the AR-process in PyStan """
27 AR_sm = stan.StanModel(file='AR1.stan')
28
29 """ And treat the previous simulations as the datasets """
30 AR_x_data = {"T": 201,
31               "x": x}
32 AR_y_data = {"T": 201,
33               "x": y}
34
35 x_fit = AR_sm.sampling(data=AR_x_data, iter=5000, warmup=1000)

```

```
36 # Increase warmup to get a better n_eff
37 y_fit = AR_sm.sampling(data=AR_y_data, iter=5000, warmup=2500)
38
39 x_fit_df = x_fit.to_dataframe()[['phi', 'mu', 'sigma']]
40 y_fit_df = y_fit.to_dataframe()[['phi', 'mu', 'sigma']]
41
42 sea.pairplot(x_fit_df)
43 sea.pairplot(y_fit_df)
```